

Scientific Computing with Python


Mini-Project: k Nearest Neighbor of images

Background

- **As our first approach, we will develop what we call a Nearest Neighbor Classifier to classify hand written digits dataset MNIST. This method will allow us to get an idea about the basic approach to an image classification problem.**
- **Example image classification dataset: MNIST.**
 - One popular toy image classification dataset is the MNIST dataset. This dataset consists of 60,000 tiny images that are 28 pixels high and wide.
- **Each image is labeled with one of 10 classes (0-9). These 60,000 images are partitioned into a training set of 50,000 images and a test set of 10,000 images.**
- **In the image below you can see 10 random example images from each one of the 10 classes:**

Introducing the MNIST problem


- MNIST (Mixed National Institute of Standards and Technology database) is a large database of handwritten digits that is commonly used for training various image processing systems.
- It consists of images of handwritten digits like these:



- The MNIST database contains **50,000** training images, **10,000** validation images and **10,000** testing images.


Flatten the 2D image into 1D vector

- We first flatten each image into a vector of $28 \times 28 = 784$ numbers. It doesn't matter how we flatten the array, as long as we're consistent between images.
- From this perspective, the MNIST images are just a bunch of points in a 784-dimensional vector space.



Result of the Flatten Operation

- The result is that the training images is a matrix (tensor) with a shape of `[50000, 784]`.
- The first dimension is an index into the list of images and the second dimension is the index for each pixel in each image.
- Each entry in the tensor is a pixel intensity between 0 and 255, for a particular pixel in a particular image.

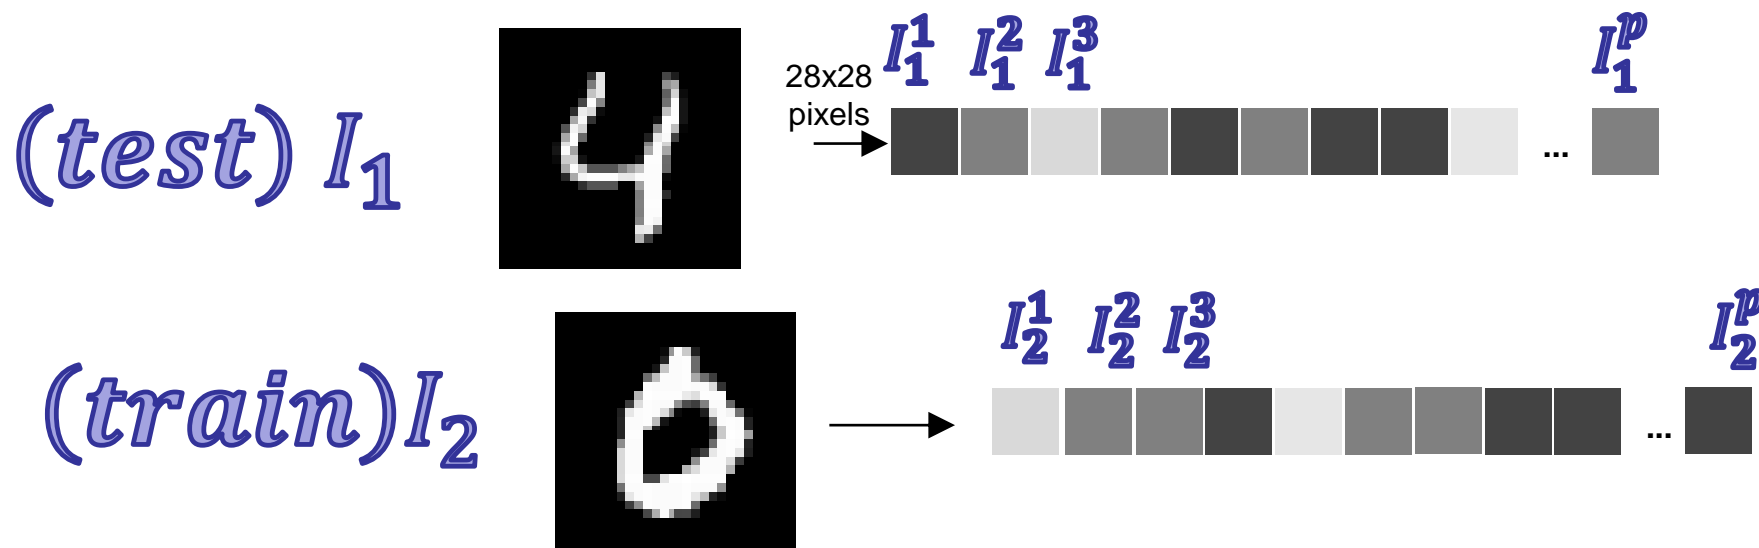


How kNN work for image classification?


- For each test image, we “**compare**” the image with all the training set images, then we will find **k** nearest neighbor images and let the k images vote for the test image and determine the label of the test image.
 - How do we compare images?

- Use the L2 distance (Euclidean distance between two vectors):

$$d_2(I_1, I_2) = \sqrt{(I_1^1 - I_2^1)^2 + (I_1^2 - I_2^2)^2 + \dots + (I_1^p - I_2^p)^2} = \sqrt{\sum_{i=1}^p (I_1^i - I_2^i)^2}$$




Visualization of K Nearest Neighbor



How do the k images vote the test image?

- For example, given the below image that we need to label, we have found 7 nearest neighbors for this image:



- So based on the 7 neighbors, 4 votes “4”, 3 votes “9”, this image will be labeled as 4