

Scientific Computing with Python

Feng Chen
HPC User Services
LSU HPC & LONI
sys-help@loni.org

Louisiana State University
Baton Rouge
May 29, 2017

Outline

- **Day 1 Basic Python Review**
- **Introducing Python Modules:**
 - ❖ Numpy
 - ❖ Scipy
- **Examples**
 - Calculate derivative of a function
 - Convert RGB image to Grayscale
 - Simple regression example
- **Exercise and Mini-Project**
 - Numpy warmup exercise
 - Calculate k nearest neighbor of image data

Day 1 Basic Python Recap

- **Python is a general-purpose interpreted, interactive, object-oriented, and high-level programming language.**
- **It was created by Guido van Rossum during 1985-1990. Like Perl, Python source code is also available under the GNU General Public License (GPL).**

Advantage of using Python

➤ Python is:

- Interpreted:
 - Python is processed at runtime by the interpreter. You do not need to compile your program before executing it. This is similar to PERL and PHP.
- Interactive:
 - You can actually sit at a Python prompt and interact with the interpreter directly to write your programs.
- Object-Oriented:
 - Python supports Object-Oriented style or technique of programming that encapsulates code within objects.
- Beginner's Language:
 - Python is a great language for the beginner-level programmers and supports the development of a wide range of applications from simple text processing to browsers to games.

Directly Run into List/Array

```
#!/usr/bin/env python
# generate array from 0-4
a = list(range(5))
print(a)
# len(a)=5
for idx in range(len(a)):
    a[idx] += 5
print(a)
```

```
[fchen14@shelob001 python]$ ./loop_array.py
[0, 1, 2, 3, 4]
[5, 6, 7, 8, 9]
```

Python Tuples

- **A Python tuple is a sequence of immutable Python objects. Creating a tuple is as simple as putting different comma-separated values.**

```
#!/usr/bin/env python
tup1 = ('physics', 'chemistry', 1997, 2000);
tup2 = (1, 2, 3, 4, 5 );
tup3 = "a", "b", "c", "d";
# The empty tuple is written as two parentheses containing nothing
tup1 = ();
# To write a tuple containing a single value you have to include a comma,
tup1 = (50,);
# Accessing Values in Tuples
print("tup1[0]: ", tup1[0])
print("tup2[1:5]: ", tup2[1:5])
# Updating Tuples, create a new tuple as follows
tup3 = tup1 + tup2;
print(tup3)
# delete tuple elements
del tup3;
print("After deleting tup3 : ")
print(tup3)
```

Scientific Computing using Python

Introducing Numpy

Numpy Overview

- **NumPy (Numeric Python) is the core package for scientific computing in Python.**
- **It is a Python library that provides a multidimensional array object, various derived objects (such as masked arrays and matrices)**
- **An assortment of routines for fast operations on arrays, including mathematical, logical, shape manipulation, sorting, selecting, I/O, discrete Fourier transforms, basic linear algebra, basic statistical operations, random simulation and much more.**
- **NumPy package provides basic routines for manipulating large arrays and matrices of numeric data.**

Basic Array Operations

- Simple array math using `np.array`
- Note that NumPy array starts its index from **0**, end at **N-1** (C-style)

To avoid module name collision inside package context

```
>>> import numpy as np
```

```
>>> a = np.array([1,2,3])
```

```
>>> b = np.array([4,5,6])
```

```
>>> a+b
```

```
array([5, 7, 9])
```

```
>>> a*b
```

```
array([ 4, 10, 18])
```

```
>>> a ** b
```

```
array([ 1, 32, 729])
```

Setting Array Element Values

```
>>> a[0]
1
>>> a[0]=11
>>> a
array([11,  2,  3,  4])
>>> a.fill(0) # set all values in the array with 0
>>> a[:]=1 # why we need to use [:]?
>>> a
array([1, 1, 1, 1])
>>> a.dtype # note that a is still int64 type !
dtype('int64')
>>> a[0]=10.6 # decimal parts are truncated, be careful!
>>> a
array([10,  1,  1,  1])
>>> a.fill(-3.7) # fill() will have the same behavior
>>> a
array([-3, -3, -3, -3])
```

Numpy Array Properties (1)

```
>>> a = np.array([0,1,2,3]) # create a from a list
# create evenly spaced values within [start, stop)
>>> a = np.arange(1,5)
>>> a
array([1, 2, 3, 4])
>>> type(a)
<type 'numpy.ndarray'>
>>> a.dtype
dtype('int64')
# Length of one array element in bytes
>>> a.itemsize
8
```

Numpy Array Properties (2)

```
# shape returns a tuple listing the length of the array
# along each dimension.
>>> a.shape # or np.shape(a)
>>> a.size  # or np.size(a), return the total number of elements
4
# return the number of bytes used by the data portion of the array
>>> a.nbytes
32
# return the number of dimensions of the array
>>> a.ndim
1
```

Numpy Array Creation Functions (1)

Nearly identical to Python's range(). Creates an array of values in the range [start,stop) with the specified step value. Allows non-integer values for start, stop, and step. Default dtype is derived from the start, stop, and step values.

```
>>> np.arange(4)
```

```
array([0, 1, 2, 3])
```

```
>>> np.arange(0, 2*np.pi, np.pi/4)
```

```
array([ 0.,  0.78539816,  1.57079633,  2.35619449,  3.14159265,  3.92699082,  4.71238898,  5.49778714])
```

```
>>> np.arange(1.5,2.1,0.3)
```

```
array([ 1.5,  1.8,  2.1])
```

ONES, ZEROS

ones(shape, dtype=float64)

zeros(shape, dtype=float64)

shape is a number or sequence

specifying the dimensions of the # array. If dtype is not specified, # it defaults to float64.

```
>>> a=np.ones((2,3))
```

```
>>> a
```

```
array([[ 1.,  1.,  1.],
       [ 1.,  1.,  1.]])
```

```
>>> a.dtype
```

```
dtype('float64')
```

```
>>> a=np.zeros(3)
```

```
>>> a
```

```
array([ 0.,  0.,  0.])
```

```
>>> a.dtype
```

```
dtype('float64')
```

Numpy Array Creation Functions (2)

```
# Generate an n by n identity
# array. The default dtype is
# float64.
```

```
>>> a = np.identity(4)
>>> a
array([[ 1.,  0.,  0.,  0.],
       [ 0.,  1.,  0.,  0.],
       [ 0.,  0.,  1.,  0.],
       [ 0.,  0.,  0.,  1.]])
```

```
>>> a.dtype
dtype('float64')
>>> np.identity(4, dtype=int)
array([[1, 0, 0, 0],
       [0, 1, 0, 0],
       [0, 0, 1, 0],
       [0, 0, 0, 1]])
```

```
# empty(shape, dtype=float64,
# order='C')
```

```
>>> a = np.empty(2)
>>> a
array([ 0.,  0.])
# fill array with 5.0
>>> a.fill(5.0)
>>> a
array([ 5.,  5.])
# alternative approach
# (slightly slower)
>>> a[:] = 4.0
>>> a
array([ 4.,  4.]])
```

Numpy Array Creation Functions (3)

```
# Generate N evenly spaced elements between (and including)  
# start and stop values.
```

```
>>> np.linspace(0,1,5)  
array([ 0. ,  0.25,  0.5 ,  0.75,  1.  ])
```

```
# Generate N evenly spaced elements on a log scale between  
# base**start and base**stop (default base=10).
```

```
>>> np.logspace(0,1,5)  
array([ 1., 1.77827941, 3.16227766,  5.62341325, 10.] )
```

Array from/to ASCII files

➤ **Useful tool for generating array from txt file**

- loadtxt
- genfromtxt

➤ **Consider the following example:**

```
# data.txt
Index
Brain Weight
Body Weight
#here is the training set
1      3.385      44.500 abjkh
2      0.480      33.38  bc_00asdk
...
#here is the cross validation set
6      27.660     115.000 rk
7      14.830      98.200 fff
...
9      4.190      58.000 kij
```


Using loadtxt and genfromtxt

```
>>> a= np.loadtxt('data.txt',skiprows=16,usecols={0,1,2},dtype=None,comments="#")
```

```
>>> a
```

```
array([[ 1.    ,  3.385,  44.5  ],
       [ 2.    ,  0.48 ,  33.38 ],
       [ 3.    ,  1.35 ,   8.1  ],
       [ 4.    , 465.    , 423.   ],
       [ 5.    ,  36.33 , 119.5  ],
       [ 6.    ,  27.66 ,  115.   ],
       [ 7.    ,  14.83 ,   98.2  ],
       [ 8.    ,   1.04 ,    5.5  ],
       [ 9.    ,   4.19 ,   58.   ]])
```

np.genfromtxt can guess the actual type of your columns by using dtype=None

```
>>> a= np.genfromtxt('data.txt',skip_header=16,dtype=None)
```

```
>>> a
```

```
array([(1, 3.385, 44.5, 'abjkh'), (2, 0.48, 33.38, 'bc_00asdk'),
       (3, 1.35, 8.1, 'fb'), (4, 465.0, 423.0, 'cer'),
       (5, 36.33, 119.5, 'rg'), (6, 27.66, 115.0, 'rk'),
       (7, 14.83, 98.2, 'fff'), (8, 1.04, 5.5, 'zxs'),
       (9, 4.19, 58.0, 'kij')],
      dtype=[('f0', '<i8'), ('f1', '<f8'), ('f2', '<f8'), ('f3', 'S9')])
```

Reshaping arrays

```
>>> a = np.arange(6)
>>> a
array([0, 1, 2, 3, 4, 5])
>>> a.shape
(6,)
>>> a.shape = (2,3) # reshape array to 2x3
>>> a
array([[0, 1, 2],
       [3, 4, 5]])
>>> a.reshape(3,2) # reshape array to 3x2
array([[0, 1],
       [2, 3],
       [4, 5]])
>>> a.reshape(2,5) # cannot change the number of elements in the array
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: total size of new array must be unchanged
>>> a.reshape(2,-1) # numpy determines the last dimension
```

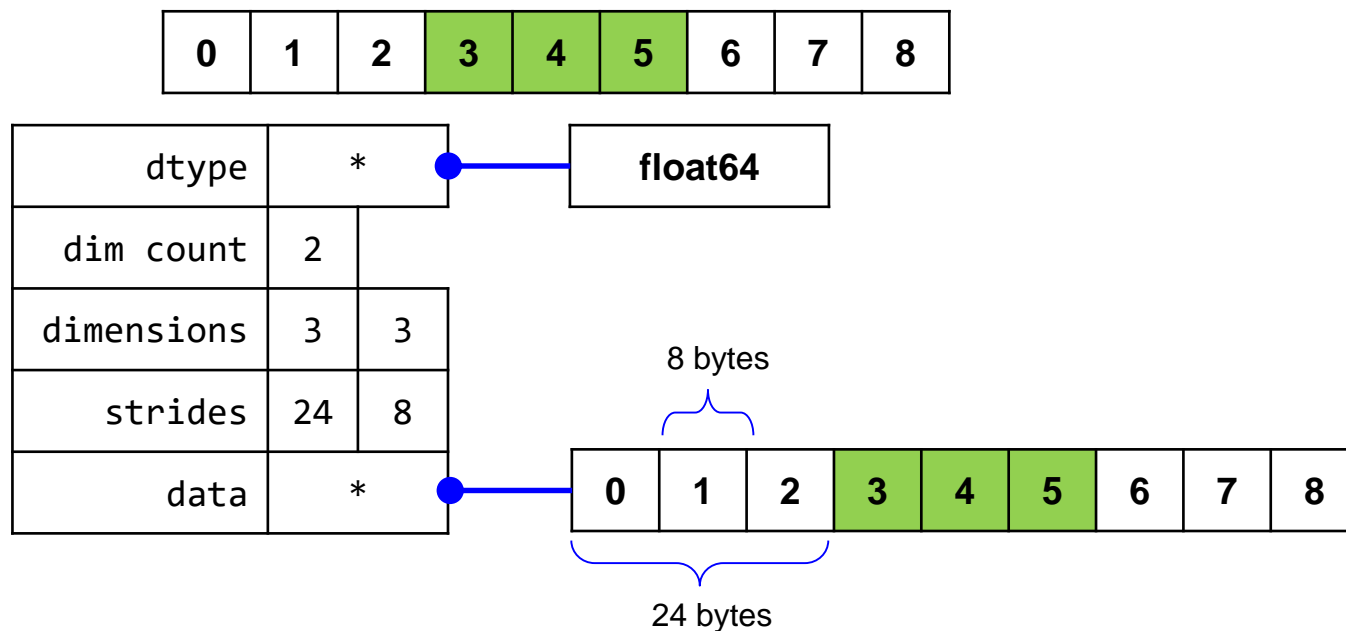
Numpy Array Data Structure

➤ Numpy view of 2D array

```
>>> a=arange(9).reshape(3,-1)
>>> a.strides
(24, 8)
>>> a.ndim
2
```

0	1	2
4	5	6
7	8	9

➤ Memory block of the 2D array



Flattening Multi-dimensional Arrays

```
# Note the difference between
# a.flatten() and a.flat
>>> a
array([[1, 2, 3],
       [4, 5, 6]])
# a.flatten() converts a
# multidimensional array into
# a 1-D array. The new array is a
# copy of the original data.
>>> b = a.flatten()
>>> b
array([1, 2, 3, 4, 5, 6])
>>> b[0] = 7
>>> b
array([7, 2, 3, 4, 5, 6])
>>> a
array([[1, 2, 3],
       [4, 5, 6]])
```

```
# a.flat is an attribute that
# returns an iterator object that
# accesses the data in the multi-
# dimensional array data as a 1-D
# array. It references the original
# memory.
>>> a.flat
<numpy.flatiter object at 0x1421c40>
>>> a.flat[:]
array([1, 2, 3, 4, 5, 6])
>>> b = a.flat
>>> b[0] = 7
>>> a
array([[7, 2, 3],
       [4, 5, 6]])
```

(Un)raveling Multi-dimensional Arrays

```
>>> a
array([[7, 2, 3],
       [4, 5, 6]])
# ravel() is the same as flatten
# but returns a reference of the
# array if possible
>>> b = a.ravel()
>>> b
array([7, 2, 3, 4, 5, 6])
>>> b[0] = 13
>>> b
array([13, 2, 3, 4, 5, 6])
>>> a
array([[13, 2, 3],
       [4, 5, 6]])
```

```
>>> at = a.transpose()
>>> at
array([[13, 4],
       [ 2, 5],
       [ 3, 6]])
>>> b = at.ravel()
>>> b
array([13, 4, 2, 5, 3, 6])
>>> b[0]=19
>>> b
array([19, 4, 2, 5, 3, 6])
>>> a
array([[13, 2, 3],
       [4, 5, 6]])
```

Four Tools in Numpy

➤ Removing loops using NumPy

- 1) Ufunc (Universal Function)
- 2) Aggregation
- 3) Broadcasting
- 4) Slicing, masking and fancy indexing

Numpy's Universal Functions

- **Numpy's universal function (or ufunc for short) is a function that operates on ndarrays in an element-by-element fashion**
- **Ufunc is a “vectorized” wrapper for a function that takes a fixed number of scalar inputs and produces a fixed number of scalar outputs.**
 - Vectorization (simplified): is the process of rewriting a loop so that instead of processing a single element of an array N times, it processes (say) 4 elements of the array simultaneously $N/4$ times.
- **Many of the built-in functions are implemented in compiled C code.**
 - They can be much faster than the code on the Python level

Ufunc Is Very Fast!

➤ Loop version

```
a=list(range(100000))
timeit [val+5 for val in a]
100 loops, best of 3: 4.94 ms per loop
```

➤ Ufunc version

```
a=np.array(a)
timeit a+5
10000 loops, best of 3: 98 µs per loop
```

➤ Speedup

- 4.94 ms / 98 µs=50 !

Ufunc: Math Functions on Numpy Arrays

```
>>> x = np.arange(5.)
>>> x
array([ 0.,  1.,  2.,  3.,  4.])
>>> c = np.pi
>>> x *= c
array([ 0.,          3.14159265,  6.28318531,  9.42477796,
        12.56637061])
>>> y = np.sin(x)
>>> y
array([ 0.00000000e+00,  1.22464680e-16, -2.44929360e-16,
        3.67394040e-16, -4.89858720e-16])
>>> import math
>>> y = math.sin(x) # must use np.sin to perform array math
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: only length-1 arrays can be converted to Python scalars
```

Ufunc: Many ufuncs available

- **Arithmetic Operators:** `+` `-` `*` `/` `//` `%` `**`
- **Bitwise Operators:** `&` `|` `~` `^` `>>` `<<`
- **Comparison Oper's:** `<` `>` `<=` `>=` `==` `!=`
- **Trig Family:** `np.sin`, `np.cos`, `np.tan` ...
- **Exponential Family:** `np.exp`, `np.log`, `np.log10` ...
- **Special Functions:** `scipy.special.*`
- ... and many, many more.

Aggregation Functions


- **Aggregations are functions which summarize the values in an array (e.g. min, max, sum, mean, etc.)**
- **Numpy aggregations are much faster than Python built-in functions**

Numpy Aggregation - Array Calculation

```
>>> a=np.arange(6).reshape(2,-1)
>>> a
array([[0, 1, 2],
       [3, 4, 5]])
# by default a.sum() adds up all values
>>> a.sum()
15
# same result, functional form
>>> np.sum(a)
15
# note this is not numpy's sum!
>>> sum(a)
array([3, 5, 7])
# not numpy's sum either!
>>> sum(a,axis=0)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: sum() takes no keyword
arguments
```

```
# sum along different axis
>>> np.sum(a,axis=0)
array([3, 5, 7])
>>> np.sum(a,axis=1)
array([ 3, 12])
>>> np.sum(a,axis=-1)
array([ 3, 12])
# product along different axis
>>> np.prod(a,axis=0)
array([ 0,  4, 10])
>>> a.prod(axis=1)
array([ 0, 60])
```

The **axes** of an array describe the order of indexing into the array, e.g., axis=0 refers to the first index coordinate, axis=1 the second, etc.



Numpy Aggregation – Statistical Methods

```
>>> np.set_printoptions(precision=4)
# generate 2x3 random float array
>>> a=np.random.random(6).reshape(2,3)
>>> a
array([[ 0.7639,  0.6408,  0.9969],
       [ 0.5546,  0.5764,  0.1712]])
>>> a.mean(axis=0)
array([ 0.6592,  0.6086,  0.5841])
>>> a.mean()
0.61730865425015347
>>> np.mean(a)
0.61730865425015347
# average can use weights
>>> np.average(a,weights=[1,2,3],axis=1)
array([ 0.8394,  0.3702])
# standard deviation
>>> a.std(axis=0)
array([ 0.1046,  0.0322,  0.4129])
```

```
# variance
>>> np.var(a, axis=1)
array([ 0.0218,  0.0346])
>>> a.min()
0.17118969968007625
>>> np.max(a)
0.99691892655137737
# find index of the minimum
>>> a.argmin(axis=0)
array([1, 1, 1])
>>> np.argmax(a,axis=1)
array([2, 1])
# this will return flattened index
>>> np.argmin(a)
5
>>> a.argmax()
2
```

Numpy's Aggregation - Summary

➤ **All have the same call style.**

- `np.min()` `np.max()` `np.sum()` `np.prod()`
- `np.argsort()`
- `np.mean()` `np.std()` `np.var()` `np.any()`
- `np.all()` `np.median()` `np.percentile()`
- `np.argmin()` `np.argmax()` . . .
- `np.nanmin()` `np.nanmax()` `np.nansum()`. . .

Array Broadcasting

- **Broadcasting is a set of rules by which ufuncs operate on arrays of different sizes and/or dimensions.**
- **Broadcasting allows NumPy arrays of different dimensionality to be combined in the same expression.**
- **Arrays with smaller dimension are broadcasted to match the larger arrays, without copying data.**

Broadcasting Rules

1. If array shapes differ, left-pad the smaller shape with 1s
2. If any dimension does not match, broadcast the dimension with size=1
3. If neither non-matching dimension is 1, raise an error.

`np.arange(3) + 5`

0	1	2
---	---	---

+

5

=

0	1	2
---	---	---

+

5	5	5
---	---	---

=

5	6	7
---	---	---

`np.ones((3,3)) + np.arange(3)`

1	1	1
1	1	1
1	1	1

+

0	1	2
---	---	---

=

1	1	1
1	1	1
1	1	1

+

0	1	2
0	1	2
0	1	2

=

1	2	3
1	2	3
1	2	3

`np.arange(3).reshape(3,1) + np.arange(3)`

0
1
2

+

0	1	2
---	---	---

=

0	0	0
1	1	1
2	2	2

+

0	1	2
0	1	2
0	1	2

=

0	1	2
1	2	3
2	3	4

Broadcasting Rules – 1D array

`np.arange(3) + 5`

0	1	2
---	---	---

+

5

1. If array shapes differ, left-pad the smaller shape with 1s

- 1) `shape=(3,)` `shape=()`
- 2) `shape=(3,)` `shape=(1,)`
- 3) `shape=(3,)` `shape=(3,)`
- 4) `final shape=(3,)`

2. If any dimension does not match, broadcast the dimension with size=1

3. If neither non-matching dimension is 1, raise an error.

0	1	2
---	---	---

+

5	5	5
---	---	---

=

5	6	7
---	---	---

Broadcasting Rules – 2D array (1)

`np.ones((3,3)) + np.arange(3)`

1	1	1		0	1	2		1	2	3
1	1	1	+	0	1	2	=	1	2	3
1	1	1		0	1	2		1	2	3

- 1) shape=(3,3) shape=(3,)
- 2) shape=(3,3) shape=(1,3)
- 3) shape=(3,3) shape=(3,3)
- final shape=(3,3)

1. If array shapes differ, left-pad the smaller shape with 1s
2. If any dimension does not match, broadcast the dimension with size=1
3. If neither non-matching dimension is 1, raise an error.

Broadcasting Rules – 2D array (2)

`np.arange(3).reshape(3,1) + np.arange(3)`

0	0	0		0	1	2		0	1	2
1	1	1	+	0	1	2	=	1	2	3
2	2	2		0	1	2		2	3	4

- 1) shape=(3,1) shape=(3)
- 2) shape=(3,1) shape=(1,3)
- 3) shape=(3,3) shape=(3,3)
- final shape=(3,3)

1. If array shapes differ, left-pad the smaller shape with 1s
2. If any dimension does not match, broadcast the dimension with size=1
3. If neither non-matching dimension is 1, raise an error.

Broadcasting Rules – Error

- The trailing axes of either arrays must be 1 or both must have the same size for broadcasting to occur. Otherwise, a **"ValueError: operands could not be broadcast together with shapes"** exception is thrown.

```
>>> a=np.arange(6).reshape(3,2)
```

```
>>> a
```

```
array([[0, 1],  
       [2, 3],  
       [4, 5]])
```

```
>>> b=np.arange(3)
```

```
>>> b
```


```
array([0, 1, 2])
```

```
>>> a+b
```

```
Traceback (most recent call last):
```

```
File "<stdin>", line 1, in <module>
```

```
ValueError: operands could not be broadcast together with shapes (3,2) (3,)
```



Slicing, Masking and Fancy Indexing

➤ **See next few slides...**

Array Slicing (1)

- `arr[lower:upper:step]`
- **Extracts a portion of a sequence by specifying a lower and upper bound. The lower-bound element is included, but the upper-bound element is not included. Mathematically: `[lower, upper)`. The step value specifies the stride between elements.**

```
# indices:           0  1  2  3  4
# negative indices:-5 -4 -3 -2 -1
>>> a = np.array([10,11,12,13,14])
# The following slicing results are the same
>>> a[1:3]
array([11, 12])
>>> a[1:-2]
array([11, 12])
>>> a[-4:3]
array([11, 12])
```

Array Slicing (2)

- **Omitting Indices:** omitted boundaries are assumed to be the beginning or end of the list, compare the following results

```
>>> a[:3] # first 3 elements
```

```
array([10, 11, 12])
```

```
>>> a[-2:] # last 2 elements
```

```
array([13, 14])
```

```
>>> a[1:] # from 1st element to the last
```

```
array([11, 12, 13, 14])
```

```
>>> a[:-1] # from 1st to the second to last
```

```
array([10, 11, 12, 13])
```

```
>>> a[:] # entire array
```

```
array([10, 11, 12, 13, 14])
```

```
>>> a[::2] # from 1st, every other element (even indices)
```

```
array([10, 12, 14])
```

```
>>> a[1::2] # from 2nd, every other element (odd indices)
```

```
array([11, 13])
```

Multidimensional Arrays

- A few 2D operations similar to the 1D operations shown above

```
>>> a = np.array([[ 0, 1, 2, 3],[10,11,12,13]], float)
```

```
>>> a
```

```
array([[ 0.,  1.,  2.,  3.],  
       [10., 11., 12., 13.]])
```

```
>>> a.shape # shape = (rows, columns)  
(2, 4)
```

```
>>> a.size # total elements in the array  
8
```

```
>>> a.ndim # number of dimensions  
2
```

```
>>> a[1,3] # reference a 2D array element  
13
```

```
>>> a[1,3] = -1 # set value of an array element
```

```
>>> a[1] # address second row using a single index  
array([10., 11., 12., -1.] )
```


2D Array Slicing

```
>>> a = np.arange(1,26)
>>> a = a.reshape(5,5) # generate the 2D array
>>> a[0,3:5]
array([4, 5])
>>> a[0,3:4]
array([4])
>>> a[4:,4:]
array([[25]])
>>> a[3:,3:]
array([[19, 20],
       [24, 25]])
>>> a[:,2]
array([ 3,  8, 13, 18, 23])
>>> a[2::2,::2]
array([[11, 13, 15],
       [21, 23, 25]])
```

1	2	3	4	5
6	7	8	9	10
11	12	13	14	15
16	17	18	19	20
21	22	23	24	25

Slices Are References

- Slices are references to memory in the original array
- Changing values in a slice also changes the original array !

```
>>> a = np.arange(5)
>>> a
array([0, 1, 2, 3, 4])
>>> b = a[2:4]
>>> b
array([2, 3])
>>> b[0]=7
>>> a
array([0, 1, 7, 3, 4])
```

Masking

```
>>> a=np.arange(10)
```

```
>>> a
```

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

a	0	1	2	3	4	5	6	7	8	9
---	---	---	---	---	---	---	---	---	---	---

```
# creation of mask using ufunc
```

```
>>> mask=np.abs(a-5)>2
```

mask	1	1	1	0	0	0	0	0	1	1
------	---	---	---	---	---	---	---	---	---	---

```
>>> mask
```

```
array([ True,  True,  True, False, False, False, False, False,  True,
        True], dtype=bool)
```

```
>>> a[mask]
```

mask	0	1	0	1						
------	---	---	---	---	--	--	--	--	--	--

```
array([0, 1, 2, 8, 9])
```

```
>>> mask=np.array([0,1,0,1],dtype=bool)
```

```
# manual creation of mask
```

```
>>> mask
```

```
array([False,  True, False,  True], dtype=bool)
```

```
>>> a[mask]
```

```
array([1, 3])
```

Masking and where

```
>>> a=np.arange(8)**2
>>> a
array([ 0,  1,  4,  9, 16, 25, 36, 49])
>>> mask=np.abs(a-9)>5
>>> mask
array([ True,  True, False, False,  True,  True,  True,  True],
      dtype=bool)
# find the locations in array where expression is true
>>> np.where(mask)
(array([0, 1, 4, 5, 6, 7]),)
>>> loc=np.where(mask)
>>> a[loc]
array([ 0,  1, 16, 25, 36, 49])
```

Masking in 2D

```
>>> a=np.arange(25).reshape(5,5)+10
>>> a
array([[10, 11, 12, 13, 14],
       [15, 16, 17, 18, 19],
       [20, 21, 22, 23, 24],
       [25, 26, 27, 28, 29],
       [30, 31, 32, 33, 34]])
>>> mask=np.array([0,1,1,0,1],dtype=bool)
>>> a[mask]    # on rows, same as a[mask,:]
array([[15, 16, 17, 18, 19],
       [20, 21, 22, 23, 24],
       [30, 31, 32, 33, 34]])
>>> a[:,mask] # on columns
array([[11, 12, 14],
       [16, 17, 19],
       [21, 22, 24],
       [26, 27, 29],
       [31, 32, 34]])
```

a[:,mask]				
0	1	1	0	1

a[mask]	0	10	11	12	13	14
	1	15	16	17	18	19
	1	20	21	22	23	24
	0	25	26	27	28	29
	1	30	31	32	33	34

Fancy Indexing - 1D

- NumPy offers more indexing facilities than regular Python sequences.
- In addition to indexing by integers and slices, arrays can be indexed by arrays of integers and arrays of Booleans (as seen before).

```
>>> a=np.arange(8)**2
>>> a
array([ 0,  1,  4,  9, 16, 25, 36, 49])
# indexing by position
>>> i=np.array([1,3,5,1])
>>> a[i]
array([ 1,  9, 25,  1])
>>> b=(np.arange(6)**2).reshape(2,-1)
>>> b
array([[ 0,  1,  4],
       [ 9, 16, 25]])
>>> i=[0,1,0]
>>> j=[0,2,1]
>>> b[i,j] # indexing 2D array by position
array([ 0, 25,  1])
```

Fancy Indexing - 2D

```
>>> b=(np.arange(12)**2).reshape(3,-1)
```

```
>>> b
```

```
array([[ 0,  1,  4,  9],
       [16, 25, 36, 49],
       [64, 81, 100, 121]])
```

```
>>> i=[0,2,1]
```

```
>>> j=[0,2,3]
```

```
# indexing 2D array
```

```
>>> b[i,j]
```

```
array([ 0, 100,  49])
```

```
# note the shape of the resulting array
```

```
>>> i=[[0,2],[2,1]]
```

```
>>> j=[[0,3],[3,1]]
```

```
# When an array of indices is used,
```

```
# the result has the same shape as the indices;
```

```
>>> b[i,j]
```

```
array([[ 0, 121],
       [121,  25]])
```

idx	0	1	2	3
0	0	1	4	9
1	16	25	36	49
2	64	81	100	121

Scientific Computing with Python

Change RGB Image to Grayscale

Using Numpy to Process Image

➤ RGB to Grayscale Conversion:

- Using simple average

$$V_{Gray} = (V_{Red} + V_{Green} + V_{Blue}) / 3$$

- Using weighted average (<https://en.wikipedia.org/wiki/Grayscale>)

$$V_{Gray} = 0.299V_{Red} + 0.587V_{Green} + 0.114V_{Blue}$$

➤ Loading and Displaying Images

Load Image


```
# import necessary images
import numpy as np
from scipy.misc import imread, imresize
import matplotlib.pyplot as plt
# To load an image, we use imread method from scipy's misc modules:
img = imread('cat.jpg')
print img.shape
```

Shape of the loaded image in ipython:

In [2]: `imread('cat.jpg')`

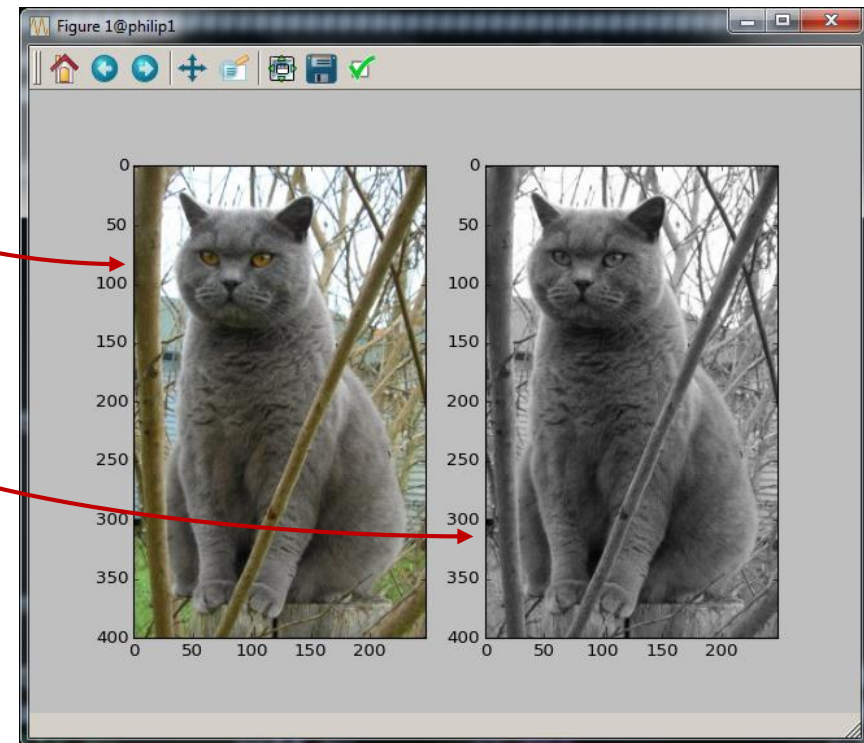
Out[2]:

```
array([[[132, 128, 117],
        [155, 151, 139],
        [181, 175, 161],
        ...,
        [ 91,  76,  57],
        [ 89,  74,  55],
        [ 86,  71,  50]]], dtype=uint8)
```



Averaging The RGB Channel Values


```
# This is simple average along axis=2
# img_tinted = np.average(img,axis=2)
# This is weighted average along axis=2
img_tinted = np.average(img,weights=[0.299,0.587,0.114],axis=2)
print img_tinted.shape
print img.shape
# plot the original image on the left
plt.subplot(1, 2, 1)
plt.imshow(img)
# plot the grayscale image on the left
plt.subplot(1, 2, 2)
plt.imshow(np.uint8(img_tinted),
            cmap='gray')
plt.show()
```



Scientific Computing with Python

Calculate Derivative

Problem Description



➤ Numerical Derivative and Integration:

$$y' = \frac{dy}{dx} \approx \frac{\Delta y}{\Delta x} = \frac{y_{i+1} - y_i}{x_{i+1} - x_i}$$

$$\int_a^b f(x) dx = \sum_{i=1}^N \frac{1}{2} (y_i + y_{i+1}) \cdot \Delta x$$

y[1]	y[2]	y[3]	y[4]
------	------	------	------

—

y[0]	y[1]	y[2]	y[3]
------	------	------	------

=

y[1]-y[0]	y[2]-y[1]	y[3]-y[2]	y[4]-y[3]
-----------	-----------	-----------	-----------

➤ How to get a vector of Δy and Δx ?

Calculate Derivative - Solution

➤ Using Numpy slicing:

```
import numpy as np
import matplotlib.pyplot as plt

# calculate the sin() function on evenly spaced data.
x = np.linspace(0, 2*np.pi, 101)
y = np.sin(x)

# use slicing to get dy and dx
dy = y[1:] - y[:-1]
dx = x[1:] - x[:-1]

dy_dx = dy/dx
```

Scientific Computing with Python

Introducing Scipy

Numerical Methods with Scipy

- **Scipy package (SCientific PYthon) provides a multitude of numerical algorithms built on Numpy data structures**
- **Organized into subpackages covering different scientific computing areas**
- **A data-processing and prototyping environment almost rivaling MATLAB**

Major modules from scipy

➤ Available sub-packages include:

- constants: physical constants and conversion factors
- cluster: hierarchical clustering, vector quantization, K-means
- integrate: numerical integration routines
- interpolate: interpolation tools
- io: data input and output
- linalg: linear algebra routines
- ndimage: various functions for multi-dimensional image processing
- optimize: optimization algorithms including linear programming
- signal: signal processing tools
- sparse: sparse matrix and related algorithms
- spatial: KD-trees, nearest neighbors, distance functions
- special: special functions
- stats: statistical functions
- weave: tool for writing C/C++ code as Python multiline strings

Scipy Example: Integration

$$\int_1^3 x^2 dx = \frac{1}{3} x^3 \Big|_1^3$$

```
#!/usr/bin/env python
import scipy.integrate as integrate
import scipy.special as special
result_integ, err = integrate.quad(lambda x: x**2, 1, 3)
result_real = 1./3.*(3.**3-1**3)

print "result_real=", result_real
print "result_integ=", result_integ
```

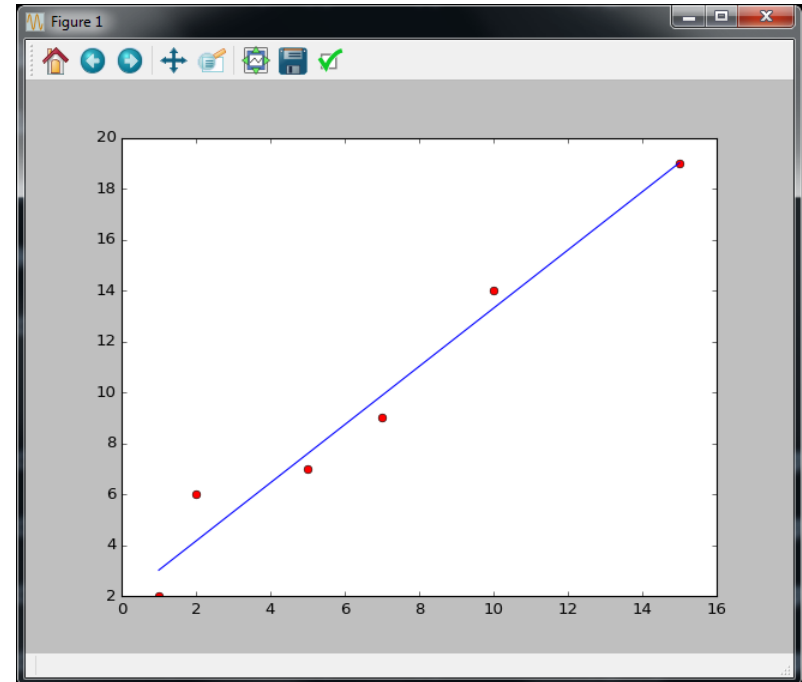
Scipy Example: Regression

```
#!/usr/bin/env python
```

```
from scipy import stats  
import numpy as np  
import matplotlib.pyplot as plt
```

```
x = np.array([1, 2, 5, 7, 10, 15])  
y = np.array([2, 6, 7, 9, 14, 19])  
slope, intercept, r_value, p_value, std_err = stats.linregress(x,y)
```

```
plt.plot(x,y,'or')  
yh = x*slope + intercept  
plt.plot(x, yh, '-b')  
plt.show()
```



Scientific Computing with Python


Mini-Project: k Nearest Neighbor of images

Background

- **As our first approach, we will develop what we call a Nearest Neighbor Classifier to classify hand written digits dataset MNIST. This method will allow us to get an idea about the basic approach to an image classification problem.**
- **Example image classification dataset: MNIST.**
 - One popular toy image classification dataset is the MNIST dataset. This dataset consists of 60,000 tiny images that are 28 pixels high and wide.
- **Each image is labeled with one of 10 classes (0-9). These 60,000 images are partitioned into a training set of 50,000 images and a test set of 10,000 images.**
- **In the image below you can see 10 random example images from each one of the 10 classes:**

Introducing the MNIST problem


- MNIST (Mixed National Institute of Standards and Technology database) is a large database of handwritten digits that is commonly used for training various image processing systems.
- It consists of images of handwritten digits like these:



- The MNIST database contains **50,000** training images, **10,000** validation images and **10,000** testing images.

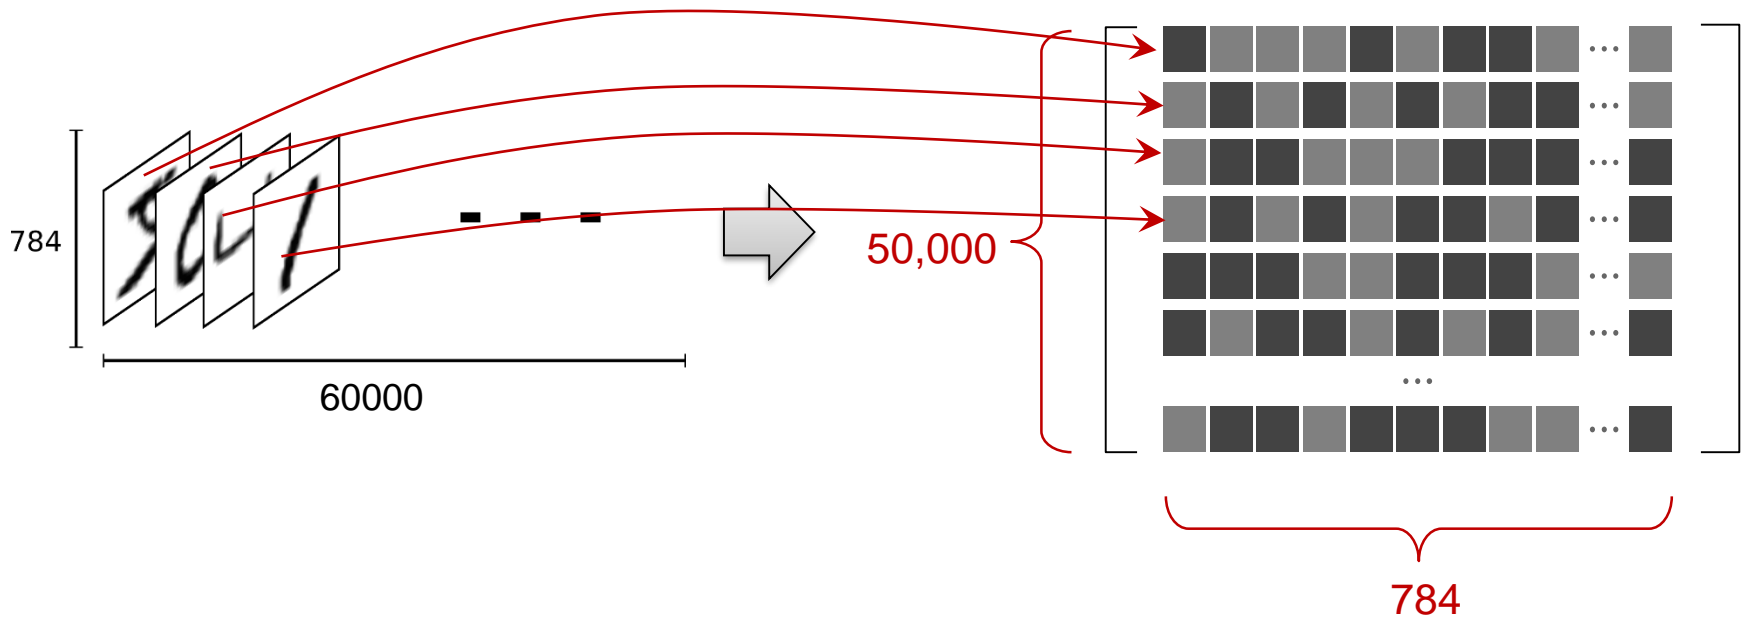
Flatten the 2D image into 1D vector

- We first flatten each image into a vector of $28 \times 28 = 784$ numbers. It doesn't matter how we flatten the array, as long as we're consistent between images.
- From this perspective, the MNIST images are just a bunch of points in a 784-dimensional vector space.



Result of the Flatten Operation

- The result is that the training images is a matrix (tensor) with a shape of **[50000, 784]**.
- The first dimension is an index into the list of images and the second dimension is the index for each pixel in each image.
- Each entry in the tensor is a pixel intensity between 0 and 255, for a particular pixel in a particular image.



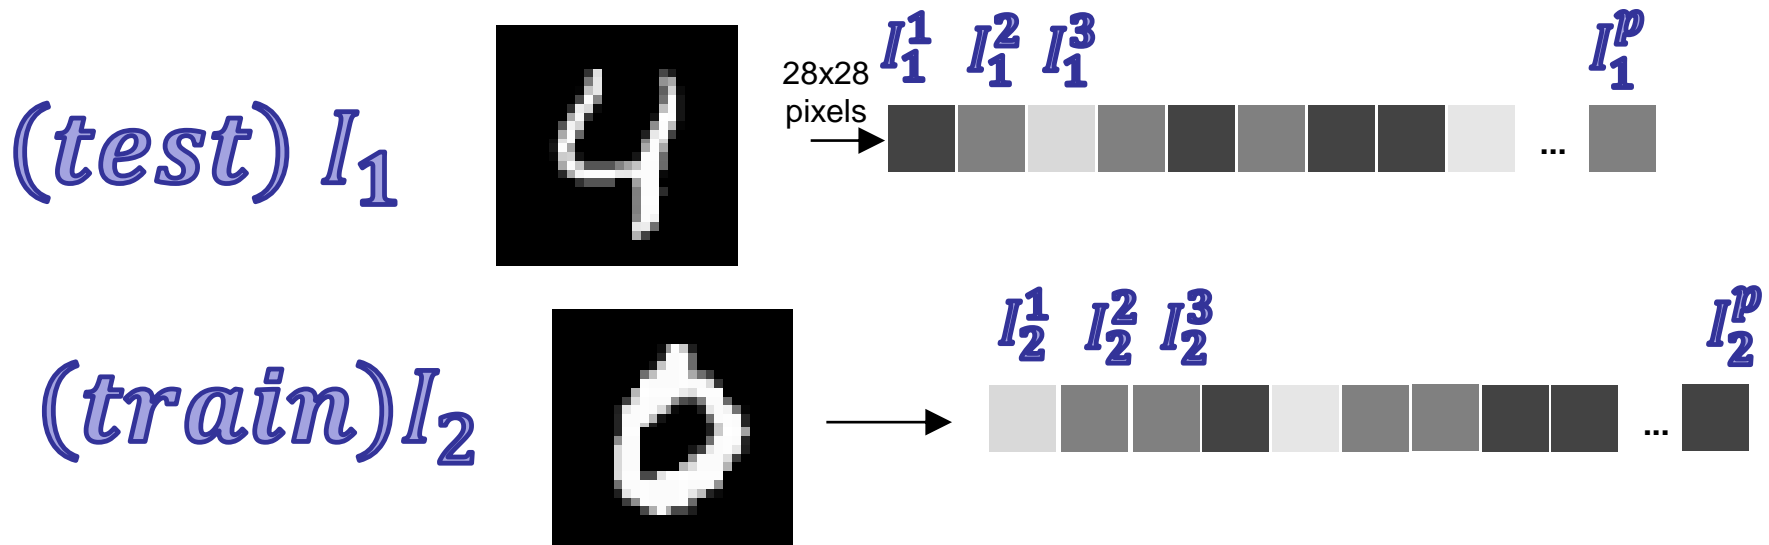
How kNN work for image classification?

- For each test image, we “**compare**” the image with all the training set images, then we will find **k** nearest neighbor images and let the k images vote for the test image and determine the label of the test image.


- How do we compare image?

- Use the L2 distance (Euclidean distance between two vectors):

$$d_2(I_1, I_2) = \sqrt{(I_1^1 - I_2^1)^2 + (I_1^2 - I_2^2)^2 + \dots + (I_1^p - I_2^p)^2} = \sqrt{\sum_{i=1}^p (I_1^i - I_2^i)^2}$$




Visualization of K Nearest Neighbor



How do the k images vote the test image?

- For example, given the below image that we need to label, we have found 7 nearest neighbors for this image:



- So based on the 7 neighbors, 4 votes “4”, 3 votes “9”, this image will be labeled as 4